COMP_SCI 496: Graduate Algorithms

Randy Truong

Spring 2024

Contents

Unit 1:
Probability + Graph Theory

Chapter 1:
A Prelude in Graph Theory and Combinatorics

This is a chapter that is devoted to Randy learning how graph theory works in finer granularity.

1. Subgraphs

Subgraph. Given a graph $G = (V, E)$, a subgraph of G, which is denoted as $G' = (V', E')$ is a graph whose vertices V' are a subset of V and whose edges E' are a subset of E.

2. Induced Subgraphs

Induced Subgraph. Given a graph $G = (V, E)$, an induced subgraph $G' = (V', E')$ is a subgraph of G that contains the following property:

- all edges $e' \in E'$ must have (both) endpoints be in V'

We define the subgraph induced by V', where V' is a set of vertices $V' \subseteq V$, as being the subgraph $G' = (V', E')$ in which all $e' \in E'$ must have both endpoints $r'_1, r'_2 \in V'$.

3. Graph Orientations

Orientation. Given an undirected graph G, an *orientation* of the graph G would be the resulting graph in which we assign each edge a direction, resulting in a directed graph.

4. Tournaments

Tournaments. Given a set of vertices V, a tournament of V, denoted as T on V, is an *orientation* of the vertices, such that the resulting graph is connected.

- Note, in a tournament, two endpoints i, j can only have a single edge between them, in which $i \rightarrow j$ or $j \rightarrow i$, but not both

5. Dominating Sets

Dominating Set. A *dominating set* of an undirected graph $G = (V, E)$ is a set $U \subseteq V$ such that every vertex $v \in V - U$ has at least one neighbor in U.

Chapter 2:
The Basic Method

---

6. The Probabilistic Method
- Powerful tool for tackling problems in discrete math
- Main idea of the method
  - ‣ **Objective.** We see kto prove that a structure with certain desired properties *exists*
  - ‣ We first define an appropriate probabilisty space of structures
  - ‣ then, we show that the desired properties hold in these structures with positive probability

7. Example 1
We note that the *Ramsey number* $R(k, \ell)$ is the smallest integer n such that in any wo-coloring of the edges of a complete graph on n vertices K_n by red and blue, either there is a red K_k (ie, a ocmplete subgraph on k vertices all of whose edges are colored red) or there is a blue K_ℓ.
- Ramsey showed that $R(k, \ell)$ is finite for any two integers k and l.
- We can obtain a lower bound for hte diagonal Ramsey numbers $R(k, k)$
  - ‣

> **Remark (Complete Graph).**
>   A *complete graph K* on n (which is denoted as K_n) is a graph in which each pair of distinct n vertices is connected together by an edge.

7.1. Prop. 1
If $\binom{n}{k} \cdot 2^{1 - \binom{k}{2}} < 1$, then $R(k, k) > n$. Thus $R(k, k) > \left\lfloor 2^{\frac{k}{2}} \right\rfloor$ for all $k \geq 3$
- Note, here $R(k, k)$ represents a graph R in which the induced monochromatic graphs must *both* be of size k, which is a stronger condition than $R(k, \ell)$

Proof. In Section 7.1, we first derive the following:
- We first seek out the probability that an induced graph K_n is monochromatic, which is equiv. to

$$\Pr[A_R] = \left(\frac{1}{2}\right)^{\binom{n}{2}} \to 2^{\binom{n}{2}} \leftrightarrow 2^{1 - \binom{n}{2}} \tag{1}$$

- We next seek the probability that any induced graph (aka a n-combination of the original k-graph) is monochromatic

$$\binom{k}{n} \cdot 2^{1 - \binom{n}{2}} \tag{2}$$

- **Observation 1.** We understand that the probability of this event occuring must be bounded by 1, thus leading us to the conclusion that

$$\binom{k}{n} \Pr[A_R] < 1 \tag{3}$$

which must imply that the probability that event $\Pr[A_R]$ doesn't occur must be non-zero

- ‣ What exactly is the *negation* of event A_R? If A_R denotes the event that the induced subgraph of K_k on R is monochromatic, then $\neg(A_R)$ must be the probability that a two-coloring of the graph R does *not* produce a monochromatic induced subgraph
- ‣ We would denote this is as the Ramsey number of induced grpah sizes n and n, or $R(n, n)$. Given that we want the negation, then we know that the size must be greater than n_0, since the Ramsey number must be n_0
- **Observation 2.** We understand that if the size of the induced graph, denoted as $n \geq 3$ and if we take the size of the graph k to be $k = \left\lfloor 2^{\frac{k}{2}} \right\rfloor$, then we know that

$$\binom{k}{n} 2^{1 - \binom{n}{2}} < \left(\frac{2^{1+\frac{n}{2}}}{n!} \right) \cdot \left(\frac{k^n}{2^{\frac{n^2}{2}}} \right) < 1 \tag{4}$$

Thus, $R(n, n) > \left\lfloor 2^{\frac{n}{2}} \right\rfloor \forall n \geq 3$

8. Optional Proof Notes
- Because there are $\binom{k}{n}$ choosings of the graph G, then it follows that at least one of these events occurring must be non-zero but strictly less than 1.
  - ‣ Thus, the inverse of this statement must be true- ($\exists$ a two-coloring of the graph G of k vertices G_k that doesn't have a monochromatic induced graph G_n)
  - ‣ This is represented as the Ramsey number $R(k, k)$, which semantically equates to smallest size of a graph that has a monochromatic edge-colored subgraph of size k
  - ‣ Given that we know that the Ramsey number refers to the smallest n for which there is a monochromatic induced subgraph of sizes k or ℓ, then it follows that in order for both monographic subgraph sto be of sizes k, then the graph must be at least the size of the graph whose Ramsey number $R(k, \ell)$ is n.

9. The Essence of the Probabilistic Method
Note the way that we evaluated this problem.
- **Objective.** Prove the existence of a good coloring K_n given a graph K.
  - ‣ We first defined what a "good" coloring was– which was a non-monochromatic graph formed from the induced graph of the two-colored graph.
  - ‣ Then, we showed that it *exists*, in a nonconstructive way
    - – We defined a probability space of events, and we narrowed that probability space down to events that described structure sof particular properties
    - – From there, we then just showed that the desired properties that we want will hold in this narrowed down probability space, with positive probability.

10. Why is this approach effective?
This approach is effective because the vast majority of probability spaces in combinatorial problems are *finite*
- Sure, we could use an algorithm to try and find such a structure with a particular property
- For example, if we wanted to actually find an edge two-coloring of K_n without a monochromatic induced graph, we could just iterate through all possible edge-colorings and find their induced grpahs.
  - ‣ Obviously, this is impractical (it's actually class $\mathbb{P}$ haha)
  - ‣ Although these problems could be solved using *exhaustive searches*, we want a faster way.
  - ‣ This is the difference between *constructivist* and *nonconstructivist* ideas in proofs
    - – Although we don't have a deterministic way of forming the graph, we are able to define an algorithm that could potentially lead to the desired graph, which, which is more effective than just trying to deterministically create one

- In the case of the Ramsey-number problem, it would be more effective to find a good coloring (a non-monochromatic induced graph) by just letting a fair coin toss decide on how to color the nodes

11. Second Look at the Probabilistic Method

12. Property S_k

We state that a tournament T has the property S_k if and only if, for every set of k Players, there is one that beats them all./

- Formally, this would mean that given a tournament $T = \langle V, E \rangle$ and subsets K of size k

$$\exists v \in T - K : (v, k) \forall k \in K \tag{5}$$

Claim. *Is it true that for every finite k that there exists a tournament T (on more than k vertices) with the property S_k?*

Proof. In order to prove this, let us consider a random tournament T.

Given this random tournament T, let's determine the probability that a node v in $T - K$ beats all of the nodes $j \in K$. This is a difficult probability to calculate, however, and it is this probability as the complement of its negation (that there isn't a node in $V - K$ that beats all the nodes $j \in K$).

Let us find probability that a fixed node v in $V - K$ beats all the nodes $j \in K$.

-

> **Remark (Tournamentk).**
>
> Because T is a tournament, we know that if we're considering a vertex v, it must be connected to all of the nodes within the subset K. Thus, there is a $\frac{1}{2}$ probability with which the edge with endpoints $v, j : j \in K$ is directed $v \to j$.

- Since there are k nodes in K and that the event of v beating a vertex j is independent, then we just find the product

$$\Pr(v \text{ beats them all}) \to \prod_{1}^{k} \left(\frac{1}{2}\right) \to \left(\frac{1}{2}\right)^{k} \leftrightarrow (2)^{-k} \tag{6}$$

From this, it follows that the probability that v doesn't beat them all is given by

$$\Pr(v \text{ does not beat them all}) =$$
$$= (1 - \Pr(v \text{ beats them all})) \tag{7}$$
$$= \left(1 - 2^{-k}\right)$$

Now, we simply just need to find the probability that *any* fixed v doesn't beat them all.

$$\text{Pr(no vertex beats them all)} =$$

$$= \text{(number of possible } v \in V - K)$$

$$\times \Pr(v \text{ doesnot beat them all})$$

$$= \prod_{v \in V - K} \Pr(v \text{ does not beat them all}) \qquad (8)$$

$$= \prod_{v \in V - K} \left(1 - 2^{-k}\right)$$

$$= \left(1 - 2^{-k}\right)^{n-k}$$

Finally, we just need to consider this scenario for all subsets K of size k in V

$$\sum_{\substack{K \subset V \\ |K|=k}} \text{Pr(no vertex beats them all)} =$$

$$= \binom{n}{k} \text{Pr(no vertex beats them all)} \qquad (9)$$

$$= \binom{n}{k} \left(1 - 2^{-k}\right)^{n-k}$$

Unit 2: Randomization Algorithms

Figure 1: how it feels to learn probability for the first time while literally in a graduate class that uses probability

Chapter 3:
Hashing (TODO)

Chapter 4:
Universal and Perfect Hashing (TODO)

Chapter 5:
Bloom Filters (TODO)

Chapter 6:
Balls and Bins (TODO)

Chapter 7:
Power of Random Choices (TODO)

Chapter 8:
Power of 2 Random Choices (TODO)

Chapter 9:
Hypercubes + Permutation Routing (TODO)

Unit 3:
Streaming Algorithms

Figure 2:

Chapter 10:
Motivation: Finding the Most Frequent Elements (TODO)

Chapter 11:
Misra-Gries Algorithm (TODO)

Chapter 12:
Count-Min Sketch (TODO)

Chapter 13:
Counting Distinct Elements (TODO)

Unit 4:
Online Algorithms

Figure 3: balling

Chapter 14:
Online Algorithms (Part 1)

---

13. Motivation
- In the former lectures, we utlized and proved the runtime and correctness of algorithms that approximated the solutions of computationally difficult problems
  - ▸ However, we now move on to a new problem:

...how can we optimize the solution to a problem in which the algorithm doesn't have all of the information that it needs?...

*...by using **Online Algorithms!***

> **Definition 1 (Online Algorithms).**
>   Algorithms used in settings where data/inputs arrive *over time,* thus requiring us to *make decisions on the fly,* without knowing what's going to happen in the future.

Of course, the opposite type of algorithm is exactly what we're used to.
- Using bubble sort? You usually know all of the inputs beforehand
- We can think of data structures themselves as being inherently online algorithms, since they handle sequences of requests, without knowledge of the future

14. List of Classical Online Algorithm Problems
- Rent or buy? (Ski rental problem)
- The elevator problem

15. The Ski Rental Problem (Rent or Buy?)
- This is the problem statement:

"Say you are just starting to go skiing. You can either rent skis for $50 or buy them again for $500. You don't know if you're going to enjoy skiing, so you opt to rent. Then you decide to go again, and again, and after a while, you realize that you have shelled out more money renting and wish you had bought right at the start."

- The intuitive solution:
  - ▸ In order to optimally solve this problem, we just need to set a

threshold. If we plan on going to ski for more than 10 days, then just buy the skis. Otherwise, if you don't plan on skiing for more than 10 days, then just rent them.
- **Objective.** What is the optimal strategy for saving money, assuming that you didn't know how often you were going to ski?

16. What exactly makes an online algorithm solution (or any solution in this case) *good?*

- We consider the *competitive ratio* of the algorithm.

> **Definition 2 (Competitive Ratio).**
> The *competitive ratio* of an online algorithm defines the **worst- case** ratio of the *online algorithm ALG* would result in on an input sequence I to the cost of the optimal offline algorithm *OPT*.
>
> $$\text{competitive ratio} = \frac{ALG(I)}{OPT(I)} \tag{10}$$

Example (Competitive Ratio 1).

- **Case 1 (Buy instantly).** Let us consider the algorithm *ALG* in which we just buy the skis instantly. We know that the worst case scenario for this problem occurs if we only ski once, ie $I = \{1\}$. If we utilize the competitive ratio definition, we observe

$$\frac{ALG(I)}{OPT(I)} = \frac{500}{50} = 10. \tag{11}$$

Equation 11 semantically equates to "the online algorithm *ALG* is ten times worse than the most optimal offline algorithm *OPT*."

Example (Competitive Ratio 2).

- **Case 2 (Rent forever).** Let us consider the algorithm *ALG* in which we just continuously rent. The worst case set of futures I is the same, as it is possible we just ski for one time. This results in

$$\frac{ALG(I)}{OPT(I)} = \frac{50 \times \infty}{50} = \infty. \tag{12}$$

This ratio is unbounded, so we can safely say that this particular *ALG* is not very good lol

Example (Competitive Ratio 3).

Let us now consider a decent strategy, which we call the *better-late-than-never* algorithm. This algorithm, denoted as *ALG*, simply states that we just rent the skis until we realize that we should have just bought them. At that point, we just buy the skis.

- Formally, if rental cost is r and the purchase cost is p, then the algorithm is to rent $\lceil \frac{p}{r} \rceil - 1$ times, then buy.

> **Theorem (Better-Late-Than-Never Performance).**
> The algorithm *better-late-than-never* has a competitive ratio ≤ 2. If the purchase cost p is an integer multiple of the rental cost r, then the competitive ratio is $2 - \left(\frac{r}{p}\right)$

Proof of BLTN Performance.

We can prove this theokem directly via case analysis.

- **Case 1.** If you went skiing less than $\lceil \frac{p}{r} \rceil$ times, which for $p = 500$ and $q = 50$ would be less than 10 times, then you are performing optimally

▸ This is because all you can do is rent, otherwise if you buy, then you're wasting money!
- **Case 2.** If you go skiing $\geq \lceil \frac{p}{r} \rceil$ times, then the best solution otherwise would have just been to buy the skis from the start or $OPT = p$. We find that in the circumstance that we do buy, ALG will pay $\left(r \times \left(\lceil \frac{p}{r} \rceil - 1 \right) \right) + p$ which is essentially equivalent to the rental price r multiplied by the maximum number of times we rent without exceeding p, plus the p we pay whenever we buy the skis. We know that arithmetically $r \times \left(\lceil \frac{p}{r} \rceil - 1 \right) + p$ must be less than or equal to $2p$. If p is a multiple of r, then it will just be $r \times \left(\left(\frac{p}{r} \right) - 1 \right) + p$, which is equal to $p - r + p \Rightarrow 2p - r$

For case 1, we demonstrated that the competitive ratio was 1. In the second case, we demonstrated that the competitive ratio was ≤ 2. Given that the worst case is 2, this must be the competitive ratio.

∎

> **Theorem (BLTN Optimality).**
> Algorithm BLTN has the best possible competitive ratio for the ski-rental problem for determiinstic algorithms when p is a multiple of r.

Proof of BLTN Optimality.
Let us consider the event that the day that you purchase the skis is the last day that you even use them.
- Consider first that this is **feasible**, since
  1. ALG never purchases $\Rightarrow$ competitive ratio is unbounded, which is undesired
  2. ALG is *deterministic*, which implies that a purchase *must* occur at some point.

Now that we've established that it *is* possible that we may purchase the skis and never use them again, let us consider the cases in which we rent *more* times than BLTN as well as *less* times:
- Renting longer than BLTN $\Rightarrow r$ increases, which implies that the competitive ratio must increase, making the algorithm worse
- Renting less than BLTN $\Rightarrow$ ratio of $\frac{r}{p}$ decreases, but this must imply that both the denominator and the numerator must decrease by $k \times r$, which also increases the competitive ratio, which is still worse.

∎

> **Note.**
> i definitely need to look this over bc i need to mathematically work out why this is sound sad

Chapter 15:
Online Algorithms (Part 2)

---

17. From last time...

In the last lecture, we were introduced to online algorithms as well as the ski rental problem.

> **Remark (Online Algorithms).**
>
> Online algorithms are algorithms that attempt to solve a problem *without* knowing the entire input space. Online algorithms will optimize their answer as inputs come in, ie *on the fly*.

> **Remark (Ski Rental Problem).**
> The ski rental problem is as follows:
>
> *Say you are just starting to go skiing.*
> *You can either rent skis for $50 or buy them again for*
> *$500. You don't know if you're going to enjoy skiing, so you*
> *opt to rent. Then you decide to go again, and again, and*
> *after a while, you realize that you have shelled out more*
> *money renting and wish you had bought right at the start.*

In the previous lecture, we discussed a **deterministic** strategy for evaluating this problem.

- Our strategy, the better-late-than-never strategy, was as follows:

    ‣ If the cost of renting was about to exceed the cost of buying the skis, then we would simply just by the skis at that point.

    ‣ We proved that in such a case, the online algorithm, denoted as *ALG*, would require that the customer rents the skis for $\left\lceil \frac{p}{r} \right\rceil - 1$ times before buying the skis outright in order to guarantee the minimal loss.

- Costs of rent vs buying
    ‣ Best evaluated using *randomization*

> **Remark (k-competitivity).**
> An algorithm is α-competitive iff
> $$\mathbb{E}[ALG(I)] \leq \alpha \cdot OPT(I) \tag{13}$$

18. Algorithm

1. Pick a "random" threshold $T \in [0, B]$

> **Remark (Deterministic Online Algorithms).**
> T was $B - 1$ in deterministic variant

2. Rent for the first T days
3. Then, buy

- Picking distribution for T?
  - ‣ Evaluate it utilizing a differnetial equation

- In principal T must be a continuous random variable for large values of B
  - ‣ Assume that we can rent up until 4.75 days, then buying

19. Picking a distribution for T

$$\text{rental cost} = \sum_i = 1^n \Pr(\text{still rent on day } i)$$

$$\text{rental cost} = \sum_i = 1^n \Pr(T > i) \tag{14}$$

where n is the number of days that we ski

$$\Pr(T \leq t) = \begin{cases} 1 \text{ , if } t \geq B \\ \frac{e^{\frac{t}{b}} - 1}{e} - 1 \text{ , for } t \in [0, B] \end{cases} \tag{15}$$

If $T = B \Rightarrow \Pr(T) = 1$, and $T = 0 \Rightarrow \Pr(T) = 0$
- Let

$$density = f_T = \frac{1}{B} \cdot \frac{e^{\frac{t}{b}}}{e - 1}, t \in [0, B] \tag{16}$$

Suppose that we ski for n days, there are two cases:
1. Case 1 $(n \leq B)$.

$$OPT = n \tag{17}$$

$$\mathbb{E}[ALG] = \mathbb{E}[\text{cost of renting}] + \mathbb{E}[\text{cost of buying}] \tag{18}$$

Let us derive $\mathbb{E}[\text{cost of buying}]$

$$\mathbb{E}[\text{cost of buying}] = \Pr(T \leq n) \cdot B \tag{19}$$

We know from cases that this equals

$$= \frac{e^{\frac{n}{B}} - 1}{e - 1} \cdot B \tag{20}$$

Now let us evaluate for the cost of renting

$$\mathbb{E}[\text{cost of renting}] \tag{21}$$

We utilize the following equality

$$\mathbb{E}[\text{cost of renting}] = \mathbb{E}[\min(n, T)] \tag{22}$$

We either rent for first n steps, or if T is sufficiently small, then buy.

- However, what is $\mathbb{E}[\min(n, T)]$

$$= \mathbb{E}_T \left[\int_0^n \mathbb{I}(t \leq T) dt \right] \tag{23}$$

Conceptually, we are integrating from $[0, n]$. It happens that T may be less tha nor equal to n. The indicator function states that the value of the function $\mathbb{I}(t \leq T) \in [0, 1]$. If $t \leq T$, then this integral is just equal to T.

- From here, what else can we do? We can utilize *linearity of expec. to swap the integral and the* $\mathbb{E}$.

$$\Rightarrow \int_0^n \mathbb{E}_T[\mathbb{I}(t \leq T)] \; dt$$
$$\Rightarrow \int_0^n \Pr(T \leq T) \; dt \; = (*) \tag{24}$$

This is just analogous utilizing a summation and finding the probability across multiple days

- Now, let us substitute the value of $\Pr(t \leq T)$ as being $(1 - \Pr(T \leq t))$.

$$(*) = \int_0^n 1 - \frac{e^{\frac{t}{B}} - 1}{e - 1} \; dt$$
$$= \int_0^n \frac{e - e^{\frac{t}{B}}}{e - 1} \; dt \tag{25}$$
$$= \frac{e}{e - 1} \cdot n - \frac{e^{\frac{n}{B}} - 1}{e - 1} \cdot B$$

Lmao, study up how to integrate again idiot

Up until this point, we have found $\mathbb{E}[\text{renting}]$ and $\mathbb{E}[\text{buying}]$, thus, we assert that

$$ALG = \frac{e}{e - 1} \cdot n \approx 1.58 \cdot n \tag{26}$$

where n is the value of OPT. This already performs a lot better than the deterministic solution of the problem.

Now let us consider the second case:

- **Case 2 $> B$.** We observe that for a sufficiently large n, then we are guaranteed that both algs would have bought skis at or before time B, thus the cost won't change.
  - ▸ Thus, the costs of both algorithms, the offline and online algorithms, would have bought the skis by time B.

Thus, we surmise that the online algorithm performs better than or equal to an optimal algorithm.

But can we do better?

There are two strategies here:

- **Approach 1.** Assume there's a distribution f with which the alg picks a threshold that is better than ours. We'll examine n and thorough computation, determine that f is the best possible function.

]

- **Approach 2.** Utilize the *minimax* principle (Yao's minimax principle for online algorithms)
  - ‣ Instead of one n, look at the distribution of n's, then do the analyss on the random distribution of n's
  - ‣ Just "do things at random for various n's"
    - – If we assume that things are done at random, then it simplifies the computations

20. Caching

* Caching: All modern processors have caches L1, L2, etc.

Let us imagine that we have a cache of size k that can store k elements/pages. Whenever we store data, we store it witihn the cache. If the cache is full, however, then we just evict an element.
- What's the optimal strategy for choosing caches to read from and evict elements from?
- LRU Cache is the most famous one
  - ‣ Basic, but practical
  - ‣ Essentially, we just keep track of whenever a cache was last used, and when we have a collision, we just evict the oldest one
- **Motivation.** We want to compare LRU with the best offline algorithm

Theorem. The competitive ratio is k, thus

$$\sup_I \frac{ALG(I)}{OPT(I)} = k \tag{27}$$

In the context of caches, we measure perforamnce based solely on the number of cache misses.
   Note. If we care about randomized algorithms, then the best competitive ratio is $O(\log k)$

Objective. We seek to prove that the number of cache misses that LRU has is bounded by $k \cdot OPT(I)$.
- Again, this could be improved utilizing ranodmization

Let us consider a sequence of instances (sequences of numbers)

$$I_1, I_2, ..., I_n \tag{28}$$

Our strat: Fix an instance I. Let us consider a block i that starts and ends with a cache miss. We find that in such a case, the optimal algorithm has 2 misses.

- How much does OPT page after the cache becomes completely full?

$$\text{numMisses}(OPT) = \text{number of blocks} \tag{29}$$

Meanwhile

$$\text{numMisses}(LRU) = k \tag{30}$$

  - ‣ Analysis
    - – How many ways does LRU page?
    - – k, since if there's more zz
  - ‣ We claim that LRU has one cache miss
    - – This is because LRU guarantees that no new element will be evicted
      - • For any distinct number in an instance I, we iwll never evict it more than once per block
    - – Thus, since OPT must evict at least one element by definition of blocks
    - – and ALG will evict at least k elements, then we complete the proof ∎
  - ‣ Let us reiterate the strats here:
    - – LRU always misses after filling the cache

- The optimal strategy? Always evict the element that is least likely to be evicted per block. This, of course, works only for offline algs

21. Next TIme (Beyond Worst Case Analysis)

- If we use resource segmentation, we find that the competitive ratio of this algorithm improves. We can demonstrate that the competitive ratio can be reduced to 2.

Chapter 16:
2-Competitive Algorithm for Online Paging
(05/14/24)

---

22. Reminders
- Homework 03 is due Friday (with extension to Monday)

23. Is 2-competitive real for LRU cache?

> **Remark ().**
>   We proved that k-competitive is the best that we can do

24. Today
- Look at randomized online algorithms
  - ‣ Martingale Algorithm
  - ‣ Beyond worst-case analysis
    - – Resource augmentation

25. Resource Augmentation
- Tactic for analyzing online algorithms
  - ‣ Compare this algorithm with the optimal offline algorithm OPT
  - ‣ In order to compensate for difference, allocate more resources for ALG

26. Scenario
- OPT
  - ‣ Has k pages
- ALG
  - ‣ Has $2k$ pages (an advantage!)
  - ‣ Could be $1.5k$ or $k + 1$, etc

> **Note.**
>   Intuition here is that: Suppose that more resources $\Rightarrow$ comparable perforamnce with OPT. We can also imagine that for OPT, resource augmentation has similar perforamnce, thus the online must be comparable

27. Why Resource Augmentation?
- Allows us to bypass the worst-case scenario

> **Theorem ().**
> The number of LRU misses is $\leq 2 \times OPT(I)$ where LRU runs with $2k$ pages and OPT with k

> **Note.**
> Proof is similar to former proofs

Proof for LRU Augmentation.

Let us consider an arbitrary sequence of futures I. We partition this block.

$$\{1, 5, 11, 1, ...\} \tag{31}$$

Let us split these into $2k$ distinct pages. For any block, consider the number of misses from both OPT and ALG.

- **OPT analysis.** In the offline cache there are k pages. If its lucky, k pages are added to the cache. This leaves, by dfn of the block, k remaining pages that weren't added, resulting in a miss. Thus, there are $\geq k$ misses per block.
- **ALG analysis.** We argue that ALG makes $\leq 2k$ misses. By definition of LRU, and the old-new page system. We will never miss on the same page within a block, but since there are $2k$ pages, then we can potentially miss *every* time

  ▸

Thus, asymptotically, the ratio is 2.

$\blacksquare$

> **Note.**
> LRU Cache is a reasonable alg because it will perform twice as worse as OPT in any time
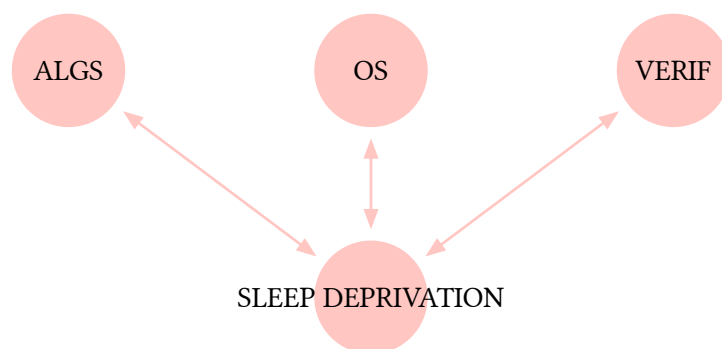
Unit 5:
FPT Algorithms

Figure 4: week 9

Chapter 17:
Preview of Second Half for Class

---

28. Motivation: Solving Computationally Hard Problems
- Imagine we had a very hard problem to solve, perhaps one that is NP-complete.
- ...What compromises can we make in order to "solve" the problem?
  - ‣ We can find the *exact* answer ↔ compromise on time
    - – Exact algorithms
  - ‣ We can *approximate* the answer ↔ save on time
    - – Approximation algorithms
  - ‣ We can *reduce* the problem as to focus on a particular parameter, rather than the entire input size ↔ save on time
    - – Fixed Parameter Tractable Algorithms

Here's kind of the itinerary of algorithm analysis for reference lol

29. Fields of Algorithm Design and Analysis

29.1. Undergraduate Algorithms
- Basically just looking at general solutions to non-computationally difficult problems
1. **Greedy Algorithms**
2. **Divide and Conquer**
3. **Dynamic Programming**

29.2. First Half of 496:
- Basically just looking at new approaches to solving these problems in order to save space and time, as well as how to optimize the solution in uncertain situations
1. **Randomized Algorithms**
2. **Streaming Algorithms**
3. **Online Algorithms**

29.3. Second Half of 496
- **How can we solve computationally-hard problems?**
1. **Fixed Parameter Tractable (FPT Algorithms).**
2. **Approximation Algorithms.**
   - **Intuition.** We can't *exactly* solve this problem in polynomial time. So let's just find an answer that is slightly worse than the optimal.
     - ‣ Similar to online algorithms, since we just care about cost
     - ‣ Emphasis on speed over correctness
3. **Exact Algorithms.**
4. **Beyond worst-case Analysis of Algorithms.**

Definitions will be revealed...

Chapter 18:
FPT Algorithms

---

30. Motivation

> **Remark (Techniques for Solving Computationally-Hard Problems).**
> One of the ways that we can solve computationally-hard problems is to utilize FPT (Fixed Parameter Tractable) algorithms. These algorithms are an alternative to **approximation algorithms**, which compromise correctness for speed. In the case of FPT, we will see that they make no compromises on correctness for speed.

- If computing the problem over an *entire* input space is too costly... what if we compute the answer utilizing smaller subproblems?

31. Motivation (cont'd): DNA Strand Length Problem

> **Note.**
> This problem isn't the *perfect* situation for utilizing FPT, since this is a polynomial time problem that we're just reducing to constant time, whereas the typical FPT problem is reducing an exponential time of n problem into an exponential time of k problem.

Let us consider the following problem:

Consider two strands of DNA. Align the two strands of DNA with the minimum possible cost (ie, with the minimum number of gaps δ).

- **Typical Solution.** We would just apply a two-dimensional dynamic programming algorithm.
    - ‣ But what if a n^2 complexity was too slow?
    - ‣ What if we knew that we just wanted to see if an alignment existed for k-cost?
- **Optimization.** Instead of calculating all alignments of all costs for DNA strands, attempt to generate the alignments to some threshold k, thus we bound the number of computations by k.

The focus of FPT algorithms is to fix some parameter k, which we can derive based on the problem's input, or as a desired result of the problem. By fixing a parameter k, we can then design an algorithm that calculates over k, thereby reducing the overall complexity, while still preserving correctness.

32. Defining FPT Algorithms

> **Definition 3 (Fixed Parameter Tractable (FPT) Algorithms).**
> FPT Algorithms are algorithms whose asymptotic running time is upper-bounded by the exponential function of a *specific parameter k* of the problem. In practice, running time is represented as this exponential function $f(k)$ multiplied by a polynomial of the input size n.
>
> $$f(k) \times poly(n) \tag{32}$$

32.1. The Inputs in More Detail

- What is k?
  - ‣ k is some parameter of the input or the solution. k is unique to the problem
- What is n?
  - ‣ The size of the input to the problem.

33. Examples of Problems that Can Be Optimized Using FPT

1. Hamiltonian Paths
2. k-paths (parameterized by k)
3. Feedback Vertex Set (parameterized by k)

34. Motivating Problem 1: Hamiltonian Paths

Problem Statement. *Suppose we have a graph G that is unweighted. Find a path that visits every vertex exactly once.*

> **Note.**
> We will get back to this one, since this one is a little difficult– first discuss the k-paths problem

35. Motivating Problem 2: k-Paths (AKA Longest Simple Path)

Problem Statement. *Given a graph G, determine if there exists a simple path of k vertices in G.*

- Evaluating the k-paths problem $\Rightarrow$ We can solve the hamiltonian path problem

35.1. Naive Solution

- **Solution 1**. Calculate all subsets of k vertices and just determine if there exists a path P between them.

- **Solution 2**. Starting from any vertex $v \in G$, just perform a breadth-first search to see if there exists a valid k-path.

- **Analysis.** Let us consider Solution (1).

  - ‣ Given a graph G such that G is Δ-regular:
    - – (suppose that $\Delta \approx n$ or $\Delta \approx \sqrt{n}$)

‣ **Runtime.** $O(\Delta^k)$, since we exhaustively search all paths of k vertices for each vertex

‣ Clearly not FPT, since $\Delta^k \approx n^k$

Given our naive solution, can we do better?

Yes! We can improve our exponential, naive algorithm by utilizing the FPT framework. We can do this by utilizing the **color coding** technique.

35.2. FPT Solution

> **Remark (k-paths Problem).**
> The graph G in the problem is *unweighted* and *undirected*.

Intuition. If we can somehow transform G into a directed graph, then we can evaulate the problem utilizing **dynamic programming**.

- Without DP, we were essentially just enumerating all possible combinations of the vertices and verifying them, which is *slow*
- By color coding the graph, we are able to filter out some combinations and make the verification process faster

> **Definition 4 (Color Coding).**
> Technique in which, given G, pick some k colors s.t., and then color all vertices $1, ..., k$ at random (independently + uniformly)

Let us now consider a basic usage of the color-coding technique, and then a more difficult one that optimizes it.

36. Evaluating k-Paths via Basic Color Coding

- This algorithm is going to inherent the "randomization" aspect of our naive algorithm by randomly selecting k-element combinations, but will now utilize **dynamic programming**, due to the added coloring

- **Algorithm.** Given a graph G:
  1. Color G using k colors (enumerated as colors $\{1, ..., k\}$)
  2. Based on this coloring, determine if a good path exists by randomly selecting k-element subsets of vertices
     ‣ ie, determine if there exists a path of colors $1 \to 2 \to ... \to (k-1) \to k$
     ‣ we can do this part using dp

> **Note.**
> This algorithm only works if the ith vertex receives the ith color.

36.1. Algorithm In Detail

- **Intuition.** Split the graph into layers based on colors $1, ..., k$
  - ‣ 1-dimensional DP
  - ‣ Utilize dynamic programming in order to demonstrate if there exists a path from 1 to k

- For each layer:
  1. Mark the ith layer (or dp_i) as reachable if and only if the vertex that precedes it is reachable

Figure 5: DP Table for the Longest Simple Path Problem

36.2. Algorithm Analysis

We derived, by definition of the algorithm, that for any combination of k-vertices, it will take k iterations to determine if its good.

- That being said, however, how many graph colorings/repetitions do we have to try in order to get the answer?

In order to evaluate this, we consider the probability that we're going find a good path.

> **Theorem (Probability of Good Path).**
> The probability that, for a path of length k, the ith vertex has the correct corresponding color i is given by
>
> $$\Pr[i\text{-th vertex gets color } i \text{ for } \{1, ..., k\}] = \left(\frac{1}{k}\right)^k = \frac{1}{k^k} \tag{33}$$
>
> at the lowest bound.

> **Note.**
> This probability is the **worst-case bound**, which occurs assuming that there is only one path. If there are multiple paths, then it's better for us.

We are able to use Equation 33 in order to determine approximately how many times we will have to run the algorithm (coloring and all) in order to obtain the answer.

> **Claim.**
> The maximum number of repetitions (of the algorithm) required to find the solution is $\approx k^k$ (ofc there's k^k different combinations).

Even though we have a bound on how many combinations we need to reveal the answer, we can concisely narrow down how many repetitions we need to guarantee that we find a path of k-colors

36.3. Making the Bound of Repetitions More Accurate

Let us now consider the failure bound of this algorithm (ie, the probability that we aren't able to find a good-colored path). In this case, we let δ be this failure bound.

> **Claim.**
> The total number of repetitions t of the Basic Graph Coloring Algorithm in order to evaluate the k-paths problem is $t \approx k^k \times \ln\left(\frac{1}{\delta}\right)$.

Proof to Claim.
- Let δ be the probability that we aren't able to find a path P of length k
- Let p be the probability that a path P of length k is good
- Let t be the number of repetitions

$$\Pr[\text{after } t \text{ attempts } \nexists \text{ path P of length } k] = (1-p)^t$$
$$\Pr[\text{after } t \text{ attempts } \nexists \text{ path P of length } k] \leq \delta$$
$$(1-p)^t \leq \delta$$
$$\ln\left((1-p)^t\right) \leq \ln(\delta)$$
$$t \times \ln(1-p) \leq \ln(\delta)$$
$$\approx t \times (-p) \leq \ln(\delta)$$
$$\approx t \geq \frac{\ln\left(\frac{1}{\delta}\right)}{p} \tag{34}$$
$$\approx t \geq \frac{\ln\left(\frac{1}{\delta}\right)}{\left(\frac{1}{k}\right)^k}$$
$$\approx t \geq \frac{\ln\left(\frac{1}{\delta}\right)}{\frac{1}{k^k}}$$
$$\approx t \geq k^k \times \ln\left(\frac{1}{\delta}\right)$$

Given Equation 34, we observe that the number of repetitions needed to guarantee that we find a path-P of length k, as required.

■

Now, we can make some claims about the true runtime of the FPT algorithm, based on the definition of FPT.

> **Claim.**
> The total runtime of the basic color coding FPT algorithm is
>
> $$k^k \times \ln\left(\frac{1}{\delta}\right) \times O(m) = $$
>
> $$= O\left(k^k\right).$$
>
> (35)

Proof.
We already know that the number of repetitions of the algorithm required in order to find a good path P of length k is given by $k^k \times \ln\left(\frac{1}{\delta}\right)$. We simply add on the new fact:

during each repetition though, the amount of time that it takes in order to verify that the path is a good path is poly(n), which we will denote as $O(m)$.

Thus, since we have to verify some $poly(n)$ paths, as well as k nodes for each path with probability $\left(\frac{1}{k}\right)$, then we know that the totla runtime must be

$$= f(k) \times poly(n)$$

$$= k^k \times \ln\left(\frac{1}{\delta}\right) \times O(m)$$

(36)

as required.

■

Although it may seem like we're done with this problem... as with all CS problems...

...how can we improve this algorithm?...

37. Evaluating k-Paths via Advanced Color Coding

> **Remark (Complexity of Naive Color Coding).**
> The asymptotic time complexity of the FPT k-Paths solution using basic color coding was $k^k \times \ln\left(\frac{1}{\delta}\right) \times O(m) = O\left(k^k\right)$.

...how can we do better...?

38. Algorithm Intuition (Advanced Color Coding)
Intuition. In our new solution, we hope to improve the runtime of the original algorithm by increasing the probability of success by *loosening* the constraint of a good path.
- In this case, we will be loosening the constraint of a good path by simply making a good path equivalent to a path in which all k vertices have distinct colors

39. Algorithm Design (Advanced Color Coding)
- Given an unweighted, undirected graph G:

- ‣ Color each vertex $v \in V$ uniquely
- ‣ Determine if $\exists$ a good path P in the resulting k-colored graph using dynamic programming
  - – A k-length path P is good if and only if all $v \in P$ have distinct colors

40. Algorithm in Detail (Advanced Color Coding)

- Given a fixed color u, find $path(u, \mathcal{S})$ where $\mathcal{S} \subseteq \{1, ..., k\}$
  - ‣ We want each path to have only one vertex of each color in $\mathcal{S}$
  - ‣ We can use S in the original problem

- The space of the DP will be $2^k \times n$ (FPT) (still way better than $n^{\log(n)}$)
  - ‣ Visit all neighbors v
  - ‣ We know that $path(u, \mathcal{S})$ exists if
    - – $\exists v \in nei(u$ where $path(v, \frac{\mathcal{S}}{color}(u)$ exists
    - – where, of course, $color(u) \in \mathcal{S}$

- We would need to repeat this algorithm $e^k \times \log\left(\frac{1}{\delta}\right)$

- Runtime is $e^k \times \log\left(\frac{1}{\delta}\right)$ multiplid by runtime of dp

41. Algorithm Analysis (Advanced Color Coding)

> **Claim.**
> The probability that all colors in a k-length path P are distinct is $\left(\frac{1}{e}\right)^k$.

Proof.
Let us first consider the probability of the event in which all colors in P are distinct, then we bound it more precisely.

We know that for any path, there are k^k possible ways to choose each vertex distinctly for a path that is k elements long. We also know that there are $k!$ permutations of k distinct elements. Thus,

$$
\begin{aligned}
\Pr[\text{all colors in } P \text{ are distinct}] &= \\
&= (\text{num possible } k \text{ color perms}) \times \Pr[\text{all elements in } P \text{ are distinct}] \\
&= k! \times \left(\frac{1}{k}\right)^k = \frac{k!}{k^k}
\end{aligned}
\tag{37}
$$

Based on Equation 37, we demonstrate that $\Pr[\text{all colors in } P \text{ are distinct}] = \frac{k!}{k^k}$. We can better bound this probability by utilizing Stirling's Inequality.

> **Remark (Stirling's Approximation).**
>
> $$
> n! \sim \sqrt{2\pi n} \times \left(\frac{n}{e}\right)^n
> \tag{38}
> $$
>
> which, we can simply just regard as
>
> $$
> n! \sim \left(\frac{n}{e}\right)^n
> \tag{39}
> $$

Using Stirling's Approximation,

$$\Pr[\text{all colors in } P \text{ are distinct}] = \frac{k!}{k^k}$$

$$\approx \frac{\left(\frac{k}{e}\right)^k}{k^k} \tag{40}$$

$$\approx \left(\frac{1}{e}\right)^k$$

Based on our calculations in Equation 40, we determine that the probability that all colors in P are distinct is $\left(\frac{1}{e}\right)^k$, as required.

∎

From here, we are able to utilize the same calculations as in Equation 36 in order to demonstrate that the runtime of the Optimized Graph Color-Coding FPT solution is $\left(\frac{1}{e}\right)^k$.

> **Claim.**
> The runtime of the Optimized Graph Color Coding FPT solution is
>
> $$e^k \times \log\left(\frac{1}{\delta}\right) \times O(m) = O(e^k). \tag{41}$$

Proof.
Utilize the same calculations as in the basic graph color coding FPT solution.

∎

Chapter 19:
Optimizing LRU Caching via Resource Augmentation (05/14/24)

42. Reminders

- Homework 03 is due Friday (with extension to Monday)

Unit 6:
Dynamic Graph Algorithms

Chapter 20:
Graph Orienting